

:: De novo Assembly ::

Introduction

This second module explores an alternative approach to reconstruct genomic data from NGS reads – **de novo Assembly**. Rather than using a mapping strategy to guide the assembly of sequence reads, de novo Assembly attempts to reconstruct genomes by exploring read overlapping and contiguity. This approach is more informative when we are dealing with a new species for which there is no reference genome or even a new strain of a well-known pathogen [14, 15].

However, the process of assembling reads into contigs can be challenging due to the huge number of reads produced from random sampling of the genome being studied. Additionally, high sequencing error rate, repeat regions/patterns and uneven sequencing depth are some of the problems that can influence the correct assembly of a genome. There are three major approaches for assembling reads: i) **overlap-and-extend**; ii) **string graph**; and iii) **de Bruijn graphs**. Overlap-and-extend methods iteratively attempt to first find read overlaps (where the suffix of a read is equal to the prefix of another read with a length that meets a defined threshold) and then extend the first read constructing a longer read (SSAKE, VCAKE and SHARCGS). The String Graph based assemblers construct a string graph for every read in which each read is a vertex and there is an edge from a vertex to another if there is read overlap (Edena and BOA). These first two approaches are more susceptible to sequencing errors and can lead to more memory consumption [15, 16].

De Bruijn Graph-based algorithms are presently the most widely used approaches and are used in several software packages for genome assembly from NGS reads (e.g. Velvet, SOAPdenovo, SPAdes, etc) [17, 18]. Graphs are mathematical structures used to model pairwise relations between objects. In essence, each vertex represents a length- k substring (k -mer) in a read and there is a directed edge from vertex u to vertex v if u and v are consecutive k -mers in a read, i.e., the last $k-1$ nucleotides of the k -mer represented by u is the same as the first $k-1$ nucleotides of the k -mer represented by v [16].

Exercise 1 – De Bruijn Graphs

Before going to the computational part we can do a simple and quick hands-on exercise that covers the basics of de Bruijn Graph based assemblies. A concept that will be

important to retain is the **concept of k-mer**. As explained above the k-mer is a substring of a read and, a read can thus have multiple k-mers depending on the k-mer length (sometimes this is also referred as hash length or word length). For example, if you have a 50bp read it can only yield one 50bp k-mer, but it can yield two 49bp k-mers [17]. The caveat of this approach is that the k-mer length you choose must always be below the read-length of your data.

Let's consider the following reads:

ATGCTA
GCTAGC
TAGCAC
CTAGCA

We can use a simple overlap and extend approach but let's use the de Bruijn Graph approach to find k-mers. List all 3 bp k-mers from the reads above:

First, link the k-mers that only differ by k-1 nucleotides. Then, you can try to establish links by creating a path between the nodes to find your contig – each k-mer is a node. The two main rules are: **visit all nodes at least once** and **use the minimal path length**.

Write the final contig here:

Exercise 2 – De novo Assembly

Now that you have attempted to do de novo assembly by hand it is time to use real data. We'll continue with the PT000033 strain from Module 1 but, instead of mapping, we will do de novo assembly. Let's go to Module2 directory by typing (from your home directory):

```
$ cd Module2
$ cp ../Module1/*trimmed* .
```

Inside this sub-directory (you can see its content using the ls command) you will find the fastq files used in Module 1, including the files produced by Trimmomatic. In this exercise we will use Velvet to assemble the PT000033 strain from sequence reads without using a reference genome [19].

You can also find in this sub-directory the reference genome used in Module 1. Although we will not use it in the assembly process, this genome file will be necessary in Exercise 3 for ordering the contigs.

This will be a simple approach to computational assembly using **Velvet**, which relies on two main programs: velveth and velvetg. velveth takes in a number of sequence files, produces a hash table, then outputs two files in an output directory (creating it if necessary), Sequences and Roadmaps, which are necessary to velvetg. The file formats supported by velveth include FASTA, FASTQ, compressed (gunzip) FASTA/FASTQ files, SAM/BAM, etc. Moreover, there are different read categories that can be specified depending on the data: -short (default), -shortPaired (for paired-end short read data), -short2 (to specify a second library), -shortPaired2 (idem), -long (for Sanger, 454 or even reference sequences) and -longPaired.

The second program, velvetg, is the core of Velvet where the de Bruijn graph is built and then manipulated.

Perhaps the most important parameter you must specify is the hash length or *k*-mer length. In fact, it cannot be inferred from the data such as other parameters that we are setting to auto or using default values. The two main rules for choosing the *k*-mer length are: it must be an odd number, to avoid palindromes (if you specify an even number Velvet will automatically decrease it); and, it must be inferior to the read length. Longer *k*-mers generally provide higher specificity but decrease coverage (and therefore sensitivity).

But, let's start assembling. Inside Module2 directory type:

```
$ velveth PT000033_41 41 -fastq -shortPaired
PT000033_1_trimmed_paired.fastq PT000033_2_trimmed_paired.fastq -
fastq -short PT000033_1_trimmed_unpaired.fastq
PT000033_2_trimmed_unpaired.fastq

# This first command creates the hash table in a new directory
PT000033_41 (it will contain the assembly using a k-mer length of
41)

$ velvetg PT000033_41 -exp_cov auto -cov_cutoff auto

# This second command produces the graph and assembles the genome.
You can find your files in the newly created PT000033_21 sub-
directory.
```

Notice that in the first command you specified two distinct libraries: a paired library present across two files using the `-shortPaired` option and two unpaired libraries using the `-short` option. We could have just used the paired library but in this way we do not lose good quality data.

Upon `velvetg` completion look at the last line summarising the results. Also, you can run the following command to find out the number of contigs and some additional statistics:

```
$ assemblathon_stats.pl ./PT000033_41/contigs.fa
```

This command will output additional statistics. The `assemblathon_stats.pl` (Keith Bradnam, UC Davis) is a Perl written script that calculates some statistics for assemblies and scaffolds (we will talk about this ahead). It was written for the Assemblathon contests to assess state-of-the-art methods in the field of genome assembly [14]. You can look at some of these metrics and take note in the table below. The output will be similar to the one below and as we are only working with assembled contigs we can focus on the third section only:

```

Number of contigs                268
Number of contigs in scaffolds    0
Number of contigs not in scaffolds 268
Total size of contigs            4343956
Longest contig                  2764650
Shortest contig                  61
Number of contigs > 1K nt        53   19.8%
Number of contigs > 10K nt       3    1.1%
Number of contigs > 100K nt      2    0.7%
Number of contigs > 1M nt        2    0.7%
Number of contigs > 10M nt       0    0.0%
Mean contig size                 16209
Median contig size               127
N50 contig length                2764650
L50 contig count                 1
contig %A                       17.25
contig %C                       32.44
contig %G                       32.91
contig %T                       17.33
contig %N                       0.06
contig %non-ACGTN               0.00
Number of contig non-ACGTN nt    0
  
```

What do these parameters mean?

- Contigs: number of contigs? Which is better, more or less contigs?
- N50: Length of the contig that contains the middle nucleotide when the contigs are ordered by size.
- Longest contig: length of the longest contig;
- Total size: Sum of all contigs lengths.

How many contigs are larger than 10 Kb? _____

Statistic	K-mer		
	41	49	55
Contigs			
Total Size			
Longest Contig			
Mean Contig Size			
N50			

Repeat this exercise using a k-mer length of 49 and 55 (**don't forget to change the directory name: PT000033_41 to PT000033_49 or PT000033_55**).

Which is the best k-mer length? _____

Upon completion of these commands you should have obtained your contigs. Each of the three assemblies are located in the respective folders. Let's look at the assembly:

```
$ cd PT000033_49

# This will get you inside the PT000033_49 subdirectory, assuming
# you are inside the Module2 directory. Recall that this is the
# assembly carried out using a k-mer length of 31bp.

# Then:

$ more contigs.fa
```

In which format are the contigs' sequences? _____

Besides the contigs file you will find other files containing the sequences, graphs and other files necessary if you want to perform a faster re-assembly of your data using different parameters.

But now that you have the contigs we can order these and/or join them in **Exercise 3**.

Optional:

Graph Visualization:

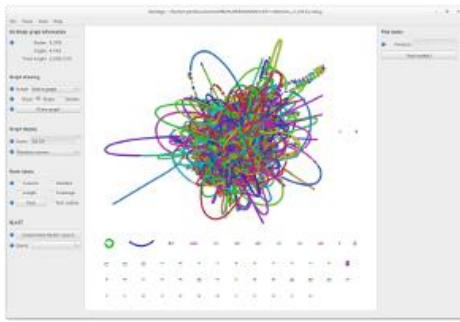
Under the assembly directories you can find the file LastGraph. It contains the final Graph of your assembly. It is possible to visualize and navigate this graph using the software Bandage. This is a program that can be installed on Windows, MacOS or Linux machines and it creates interactive visualizations of assembly graphs [20].

To open Bandage double-click on its shortcut at the desktop or, from any terminal window just type:

```
$ Bandage
```

This will start the Bandage window, go to *File > Load Graph* and select the LastGraph file from an assembly at your choice to open the Graph of your assembly. To visualize this Graph click on *Draw Graph*. Since this is a bacterial genome you can draw the entire graph. You can compare the different assemblies using this approach.

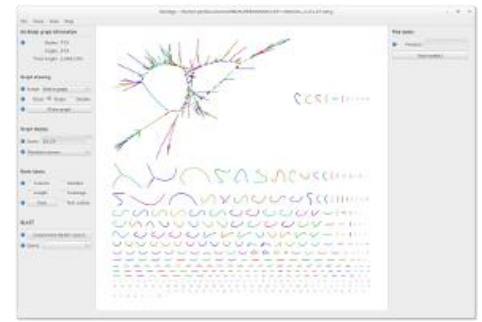
The graph for your assembly will show the different nodes and be similar to this:



K-mer of 21. This value is too small, resulting in short contigs and many connections, giving a dense tangled graph.



K-mer of 77. This is a good balance, giving long contigs that are well connected.



K-mer of 127. This value is too large, resulting in the graph breaking into many discontinuous pieces.

Other software for assembly statistics/QC:

There are more complete alternatives to evaluate the assembly quality, namely, using BUSCo or QUAST. QUAST for example compares against a reference genome, provided that one exists. Busco offers another alternative by checking the degree of representation of a set of core genes in each assembly.

You can run QUAST in this VM by:

```

$ cd PT000033_41

$ quast.py -o quast_qc -r ../NC000962_3.fasta contigs.fa

# Inside the newly created quast_qc directory look at the report
files:

$ cat quast_qc/report.tsv
    
```

For BUSCO analysis in the same directory:

```
$ conda activate busco

$ busco -i contigs.fa -l actinobacteria_phylum_odb10 -o
busco_analysis -m genome

$ conda deactivate
```

Additional notes:

The assembly exercise carried out in this exercise is a simplistic approach. In a real situation it is necessary to carefully estimate several parameters such as the expected coverage cut-off, expected coverage, k-mer length, insert size, etc. VelvetOptimiser script (<https://github.com/tseemann/VelvetOptimiser>) deals with these problems and carries out several Velvet assemblies while simultaneously adjusting for several of these parameters. You can input a range of K-mer lengths .

A different approach is implemented in SPAdes by using multi k-mer assemblies where you specify several k-mer lengths [18]. SPAdes takes as input paired-end reads, mate-pairs and single (unpaired) reads in FASTA and FASTQ. Besides SPAdes other assemblers implement a multi K-mer strategy (IDBA, Megahit, GATB-Pipeline) which always performs better than single K-mer strategies. The downside, and the main reason to why we do not used this strategy in this module, are the longer run times and the fact that these are more resource-demanding softwares.

Exercise 3 –Contig Ordering and Scaffolding

In this next exercise we will pick up from where we left in the previous exercise. Now that you have your contigs assembled from raw sequence reads we can order the contigs obtained using a suitable reference genome for this process (usually the closest high-quality finished genome). In this exercise we will use the reference genome used in Module 1 (M. tuberculosis H37Rv, GenBank Acc: NC000962.3). This reference-aided contig ordering becomes important as we progress towards discovery and comparative genomics processes. Besides contig ordering, the process of joining the contigs is called

scaffolding. Contigs within a scaffold are separated by N letters which represent gaps in the sequence. A scaffold is therefore chains of ordered contigs created using information about their relative position against a reference genome.

The software we will be using for this task at hand is called ABACAS [21]. This is a stand-alone program intended to rapidly contiguate (align, order, orientate), visualize and design primers to close gaps on shotgun assembled contigs based on a reference sequence. ABACAS is implemented in Perl and requires MUMmer and (optionally) BLAST installed on the local machine.

Let's order the contigs by simply typing the following:

```
$ cd PT000033_49
$ cp ../NC000962_3.fasta ./
$ abacas.pl -r ../NC000962_3.fasta -q contigs.fa -p promoter -b
-d -a
# Abacas will align the contigs present in the contigs file
(contigs.fa) against the reference genome and will do so based on
the promoter algorithm
```

You can choose to either do an alignment based on NUCmer (option: -p nucmer) or PROmer (option: -p promoter) algorithms. NUCmer is a Perl script pipeline for the alignment of multiple *closely related* nucleotide sequences. PROmer (PROtein MUMmer) is a close relative to the NUCmer script. It follows the exact same steps as NUCmer and even uses most of the same programs in its pipeline, with one exception - all matching and alignment routines are performed on the six-frame amino acid translation of the DNA input sequence. This provides promoter with a much higher sensitivity than nucmer because protein sequences tend to diverge much slower than their underlying DNA sequence. Therefore, on the same input sequences, promoter may find many conserved regions that nucmer will not, simply because the DNA sequence is not as highly conserved as the amino acid translation.

By the end of this ordering and scaffolding process, ABACAS will show a screen like the one below:

```
*****
* ABACAS: Algorithm Based Automatic Contiguation of Assembled Sequences *
*
*
* Copyright (C) 2008-11 The Wellcome Trust Sanger Institute, Cambridge, UK. *
* All Rights Reserved. *
*
*****

# Checking user options:
# -r Reference=./NC000962_3.fasta
# -q Query=contigs.fa
# -p promer
# -d use default setting i.e. --mumreference in promer
# -b print multifasta file of contigs in bin to file
# Input checking done!!
# PREPARING DATA FOR promer
# Ordered contigs = 374
# Bin contigs = 68
# FINISHED CONTIG ORDERING
#
# To view your results in ACT
#     Sequence file 1: NC000962_3.fasta
#     Comparison file 1: contigs.fa_NC000962_3.fasta.crunch
#     Sequence file 2: contigs.fa_NC000962_3.fasta.fasta
#
#     ACT feature file is: contigs.fa_NC000962_3.fasta.tab
#     Contigs bin file is: contigs.fa_NC000962_3.fasta.bin
#     Gaps in pseudomolecule are in: contigs.fa_NC000962_3.fasta.gaps
#-----
```

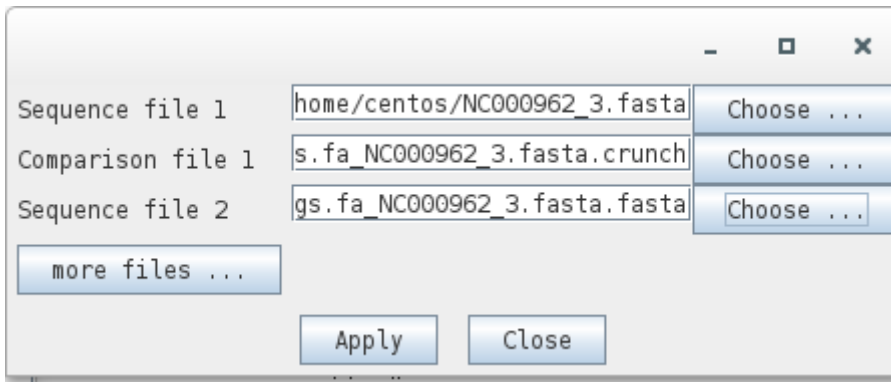
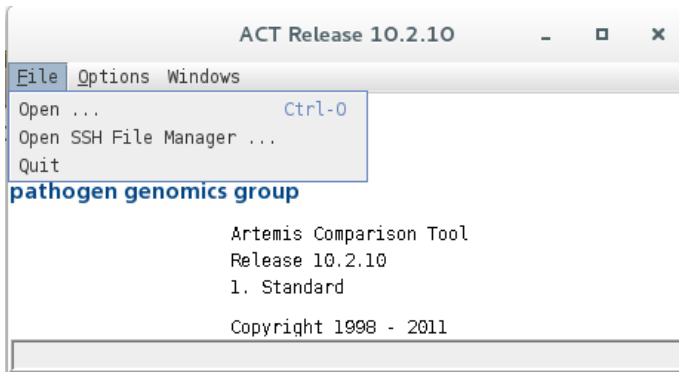
Now, what is ACT? ACT stands for Artemis Comparison Tool, a Java-written application for displaying pairwise comparisons between two or more DNA sequences (basically ACT has several Artemis windows) [22]. It can be used to identify and analyze regions of similarity and difference between genomes and to explore conservation of synteny, in the context of the entire sequences and their annotation. It can read complete EMBL, GENBANK and GFF entries or sequences in FASTA or raw format.

We can go ahead and have a look at the comparison that ABACAS already and present on the comparison file (*.crunch).

ACT can be started from any terminal window by typing:

```
$ act
```

ACT will start its main window, go to File > Open and select the files that ABACAS indicates from comparison:




Sequence 1: NC000962_3.fasta

Comparison file 1: contigs.fa_NC00096_3.fasta.crunch

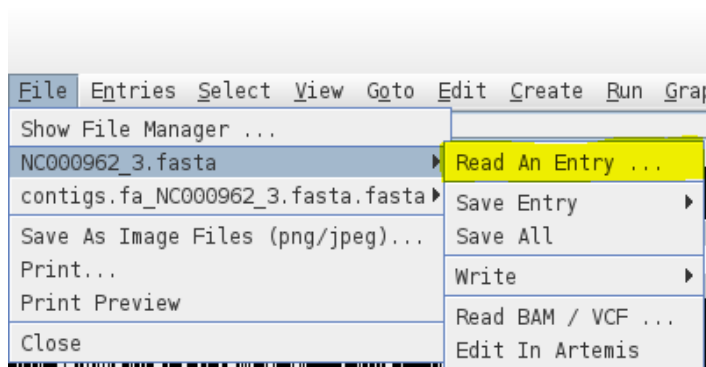
Sequence 2: contigs.fa_NC00096_3.fasta.fasta

Sequence 2 file now contains your contigs stitched to one another and separated by N's (scaffold) highlighting possible gaps in the assembly or regions not present in the reference genome. Look at the ACT comparison.




Is synteny conserved across these two genomes? What did you expect from *M. tuberculosis*?

Besides looking at the synteny of these large regions we can look at the conservation of gene synteny and distribution across this comparison. To do this we can read an entry onto of the sequences by going to *File > NC000962_3.fasta > Read an Entry...*



You can open the NC000962_3.gbk file used in the previous file. Remember this file is in the Genbank format, it contains the sequence along with genome annotation. Here, we are reading it onto the NC000962_3.fasta sequence and since the sequence is already present ACT will read the annotation only. You can see the annotation appearing on the top part of the comparison window.

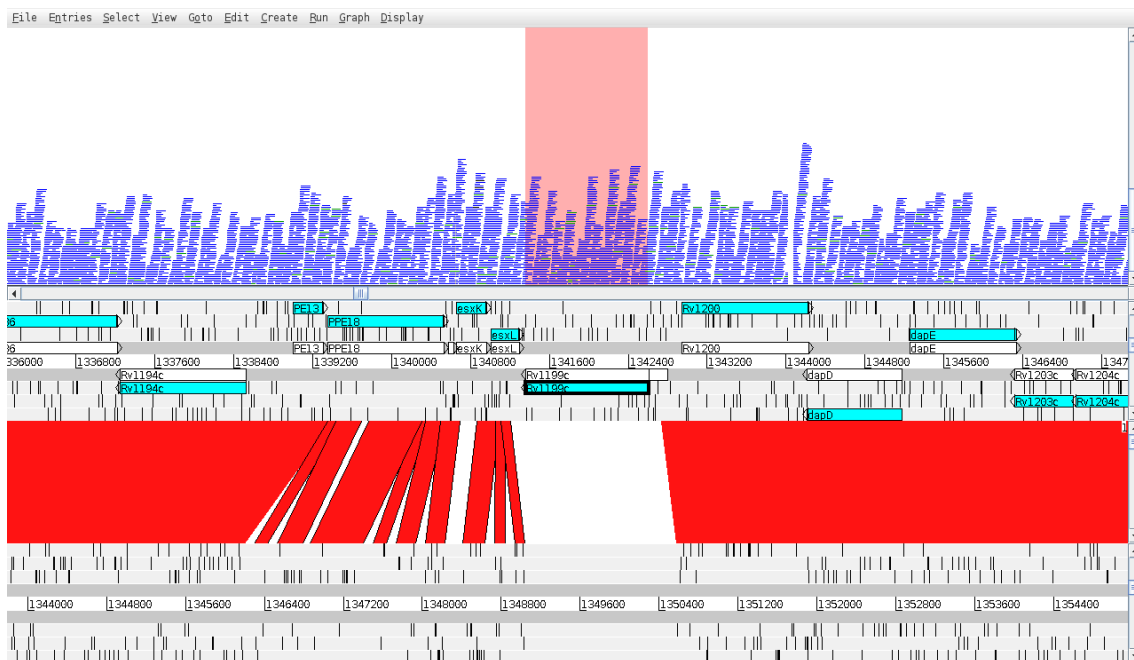


Take a look at the gaps in your assembly, where do the majority of these occur? Why is that?

We will not cover in this module how we can close these gaps but there are different strategies (see Additional notes below). We can also look at the mapped data from the previous module. Go to *File > NC000962_3.fasta > Read BAM/VCF...* and select the BAM file **PT000033.sorted.bam** in the **Module1** directory (remember that the correspondent .bai file must be in the same directory).

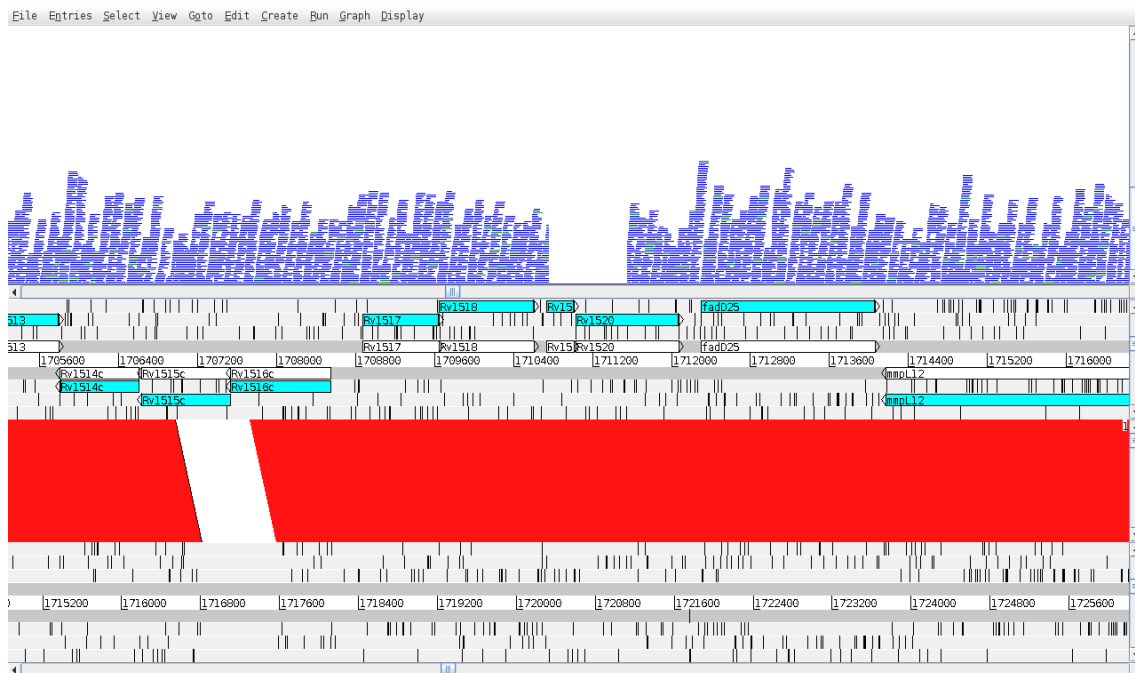
What do you see? Are most of these gapped regions real? _____

Look at the following example (around nucleotide 1341600) and locate it on your assembly:



What is happening here? _____

And here (nt 17111200):



Additional notes:

Contig ordering can also be carried out in Mauve, a genome aligner tool implemented with a Graphical User Interface (GUI), and therefore, more user friendly. The advantage of ABACAS over Mauve is that you can more easily implement contig ordering and scaffolding on a script that you run through the command-line.

For comparison purposes Mauve is also a great tool for multi-genome alignment (in particular using its Progressive Mauve algorithm). You can align your contigs against the contigs of other assembled strains, a reference strain and so on. Furthermore, during the alignment process, Mauve identifies conserved segments that appear to be internally free from genome rearrangements. Such regions are referred to as Locally Collinear Blocks (LCBs).

Also, ABACAS is part of a software bundle developed at the Wellcome Trust Sanger Institute called PAGIT. PAGIT is designed to finish high quality draft genome and, includes: ABACAS; IMAGE, an iterative approach for closing gaps in assembled genomes using mate pair information; iCORN, that enables errors in the consensus sequence to be corrected by iteratively mapping reads to the current assembly; and RATT, a tool to

transfer the annotation from a reference genome, or an earlier assembly, onto the latest assembly.

Exercise 4 - Genome Annotation

In the last exercise of this module we will attempt to find features along the scaffold created earlier and add some information concerning it to genome. This is a process called annotation and, it is necessary if we want to take full advantage of having the genome sequence and move towards functional genomics [15]. This annotation process is made with biologically relevant information that can range from gene models, functional data (including gene ontology or “Kyoto Encyclopedia of Genes and Genomes”, KEGG, pathways) to epigenetic or microRNA modifications. Generally, this procedure is limited to the annotation of protein coding sequences and may include the annotation of ribosomal and transfer RNAs.

We can take three different approaches to annotation: web-based tools (RAST, NCBI annotation); command-line tools that perform *de novo* gene discovery (Prokka and DIYA); and, transfer of annotation data from a reference genome such as the one we used in the previous exercise (RATT, BG-7) [21, 23, 24].

Moreover, annotation will imply a different format other than FASTA. While the latter only stores the sequence, it is possible to have a second file in the GFF (General Feature File) format containing the annotation. Alternatively, we can combine these data on a GenBank file. You can look up online the structure of both formats.

Prokka Annotation:

In this exercise we will take the first two approaches and we will start with annotation with Prokka, a software tool to annotate bacterial, archaeal and viral genomes quickly and produce standards-compliant output files [23]. Prokka uses a variety of databases when trying to assign function to the predicted CDS features. It takes a hierarchical approach to make it fast. The initial core databases are derived from UniProtKB.

To start the annotation let's type:

```
$ cd /home/centos/Module2/PT000033_49

$ prokka --outdir ./PT000033_prokka --prefix PT000033
contigs.fa_NC000962_3.fasta.fasta

# Alternatively, for a monomorphic species such as Mycobacterium
tuberculosis you can opt by transferring the annotation from M.
tuberculosis H37Rv reference strain using the --proteins option:

$ cp ~/course_files/NC000962_3.gbk .

$ prokka --outdir ./PT000033_prokka --prefix PT000033 --proteins
NC000962_3.gbk contigs.fa_NC000962_3.fasta.fasta
```

Here, you have carried out a simple annotation procedure using Prokka's core database. With the `--outdir` option you specified the creation of a new output directory and the `--prefix` option allow you to set the prefix name of your files. You can look at other Prokka options by simply typing:

```
$ prokka
```

Meanwhile the annotation process will take about 16 minutes, but when finished you can find the new PT000033_prokka directory. Let's go inside:

```
$ cd PT000033_prokka
```

Prokka will have outputted the following files:

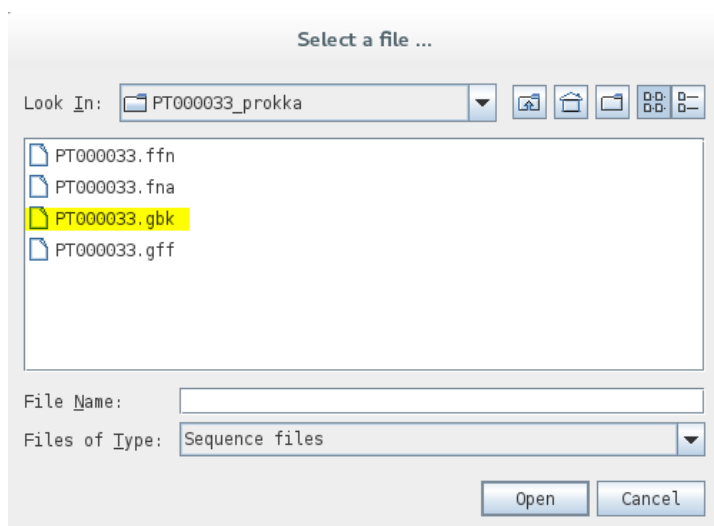
Extension	Description
.gff	This is the master annotation in GFF3 format, containing both sequences and annotations. It can be viewed directly in Artemis or IGV.
.gbk	This is a standard Genbank file derived from the master .gff. If the input to prokka was a multi-FASTA, then this will be a multi-Genbank, with one record for each sequence.
.fna	Nucleotide FASTA file of the input contig sequences.
.faa	Protein FASTA file of the translated CDS sequences.
.ffn	Nucleotide FASTA file of all the prediction transcripts (CDS, rRNA, tRNA, tmRNA, misc_RNA)
.sqn	An ASN1 format "Sequin" file for submission to Genbank. It needs to be edited to set the correct taxonomy, authors, related publication etc.

.fsa	Nucleotide FASTA file of the input contig sequences, used by "tbl2asn" to create the .sqn file. It is mostly the same as the .fna file, but with extra Sequin tags in the sequence description lines.
.tbl	Feature Table file, used by "tbl2asn" to create the .sqn file.
.err	Unacceptable annotations - the NCBI discrepancy report.
.log	Contains all the output that Prokka produced during its run. This is a record of what settings you used, even if the --quiet option was enabled.
.txt	Statistics relating to the annotated features found.
.tsv	Tab-separated file of all features: locus_tag,ftype,gene,EC_number,product

Next, open up Artemis:

```
$ art
```

And go to File > Open and open the PT000033.gbk file:



You can see that your scaffold from the previous exercise is now annotated:

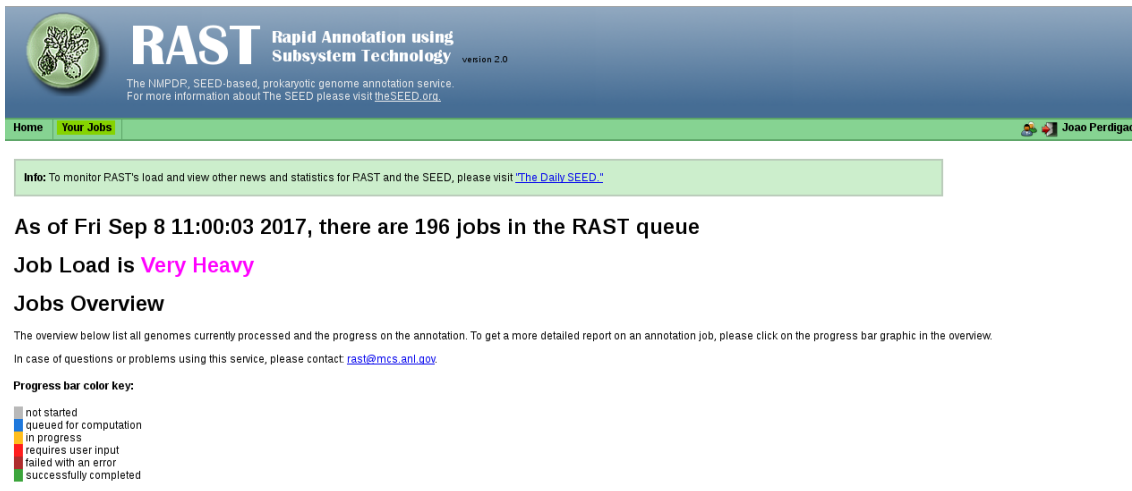


As a side note, to improve Prokka annotations it is possible to use user-define databases from closely related genomes.

RAST Annotation:

RAST (Rapid Annotation using Subsystem Technology) is a fully-automated service for annotating complete or nearly complete bacterial and archaeal genomes available at <http://rast.nmpdr.org/>. RAST is designed to rapidly call and annotate the genes of a complete or essentially complete prokaryotic genome [24]. RAST, Rapid Annotations based on Subsystem Technology, uses a "Highest Confidence First" assignment propagation strategy based on a more sophisticated database based on manually curated subsystems and subsystem-based protein families that automatically guarantees a high degree of assignment consistency. RAST returns an analysis of the genes and subsystems in your genome, as supported by comparative and other forms of evidence. Despite being an on-line tool, it is possible to use RAST on the command-line through the myRAST toolkit, therefore allowing the integration of RAST in scripts and pipelines [24].

To use RAST you need to set up a free account.



Once you have registered to RAST and your account is active, log in to your RAST homepage and go to *Your Jobs > Upload New Job*.

Upload a Genome

A prokaryotic genome in one or more contigs should be uploaded in either a single [FASTA](#) format file or in a Genbank format file. Our pipeline will use the taxonomy identifier as a handle for the genome. Therefore if at all possible please input the numeric taxonomy identifier and genus, species and strain in the following upload workflow.

Please note, that only if you submit all relevant contigs (i.e. all chromosomes, if more than one, and all plasmids) that comprise the genomic information of your organism of interest in one job, Features like *Metabolic Reconstruction* and *Scenarios* will give you a coherent picture.

If you wish to upload multiple genomes at once, you may be interested in using the batch upload interface that is available in the [myRAST distribution](#). See [this tutorial](#) for more information on this capability.

Confidentiality information: Data entered into the server will not be used for any purposes or in fact integrated into the main SEED environment, it will remain on this server for 120 days or until deleted by the submitting user.

If you use the results of this annotation in your work, please cite:

- The RAST Server: *Rapid Annotations using Subsystems Technology*. Aziz RK, Bartels D, Best AA, DeJongh M, Disz T, Edwards RA, Formsma K, Gerdes S, Glass EM, Kubal M, Meyer F, Olsen GJ, Olson R, Osterman AL, Overbeek PA, McNeil LK, Paarmann D, Paczian T, Parrello B, Pusch GD, Reich C, Stevens R, Vassileva O, Vonstein V, Wilke A, Zagnitko O. *BMC Genomics*. 2008; [PubMed entry]
- The SEED and the Rapid Annotation of microbial genomes using Subsystems Technology (RAST). Overbeek R, Olson R, Pusch GD, Olsen GJ, Davis JJ, Disz T, Edwards RA, Gerdes S, Parrello B, Shukla M, Vonstein V, Wattam AR, Xia F, Stevens R. *Nucleic Acids Res*. 2014 [PubMed entry]

File formats: You can either use [FASTA](#) or Genbank format.

- If in doubt about FASTA, [this service](#) allows conversion into FASTA format.
- Due to limits on identifier sizes imposed by some of the third-party bioinformatics tools that RAST uses, we limit the size of contig identifiers to 70 characters or fewer.
- If you use Genbank, you have the option of preserving the gene calls in the options block below. By default, genes will be recalled.

Please note: This service is intended for complete or nearly complete prokaryotic genomes, phages, or plasmids.

File Upload:
Sequences File contigs.fa_NC000962_3.fasta.fasta

On this next screen select the scaffold file(contigs.fa_NC000962_3.fasta.fasta). On the next screen you can review some sequence statistics before proceeding:

Review genome data

We have analyzed your upload and have computed the following information.

Contig statistics

Statistic	As uploaded	After splitting into scaffolds
Sequence size	4477105	4341263
Number of contigs	1	385
GC content (%)	65.4	65.4
Shortest contig size	4477105	97
Median sequence size	4477105	3013
Mean sequence size	4477105.0	11276.0
Longest contig size	4477105	98726
N50 value		30442
L50 value	1	44

Below, in this same page you will be asked to introduce some information concerning the organism. Since we are working with *Mycobacterium tuberculosis* you can enter NCBI Taxonomy Id 1773 and click on the “Fill in form based on NCBI Taxonomy-ID” button as it will fill the form automatically:

Please enter or verify the following information about this organism:

- PAST bases its genome identifiers on NCBI taxonomy-IDs.
- If you provide a valid taxonomy-ID, PAST will attempt to fill in the genome metadata for you.
- If you leave the taxonomy-ID field blank, PAST will assign a meaningless taxonomy-ID, and you will need to fill in the below genome metadata manually.
- If you plan on submitting this genome to [PATRIC](#) you will need to provide the most descriptive NCBI taxonomic grouping possible. If you leave the taxonomy-ID field blank, PAST will assign a meaningless taxonomic identifier and the genome will not be suitable for submission to PATRIC. We discuss the motivation and process for submitting your genome to PATRIC [in this document](#).
- You may search for the taxonomy-ID of your organism using the search facilities at the [NCBI taxonomy browser](#).

Genome information:

Taxonomy ID:	1773	Fill in form based on NCBI taxonomy-ID.
Taxonomy string:	Bacteria; Terrabacteria group; Actinobacteria; Actinobacteria; Corynebacteriales; Mycobacteriaceae; Mycobacterium; Mycobacterium tuberculosis complex	
Domain:	<input checked="" type="radio"/> Bacteria <input type="radio"/> Archaea <input type="radio"/> Virus	
Genus:	Mycobacterium	
Species:	tuberculosis	
Strain:		
Genetic Code:	<input checked="" type="radio"/> 11 (Archaea, most Bacteria, most Viri, and some Mitochondria) <input type="radio"/> 4 (Mycoplasmata, Spiroplasmata, Ureoplasmata, and Fungal Mitochondria)	

[Use this data and go to step 3](#)

- If you enter a valid NCBI taxonomy-ID and click “Fill in form based on NCBI taxonomy-ID,” PAST will attempt to automatically fill in the form below. You may then edit any incorrect field values before going to the next step.
- If you do not know the taxonomy-ID of your genome, please leave the taxonomy-ID field blank, and fill in the fields manually.
- If you leave this field blank, PAST will fill in a dummy taxonomy string of the form “Domain; genus species strain”, based on the form entries below.
- E.g., “*Escherichia*”. If you do not know the genus, leave blank, and it will default to “Unknown”.
- E.g., “*coli*”. If you do not know the species, leave blank, and it will default to “sp.”.
- E.g. “*str. K12 substr. MG1655*”. This field is optional. (May also be used as a comment.)

In the last step you can chose some RAST annotation parameters. We will go with the default parameters:

Upload a Genome

Complete Upload

Please consider the following options for the RAST annotation pipeline:

RAST Annotation Settings: Choose RAST annotation scheme: Classic RAST Select gene caller: RAST Select FIGfam version for this run: Release70 Automatically fix errors? <input checked="" type="checkbox"/> Yes Fix frameshifts? <input type="checkbox"/> Yes Build metabolic model? <input type="checkbox"/> Yes Backfill gaps? <input checked="" type="checkbox"/> Yes Turn on debug? <input type="checkbox"/> Yes Set verbose level: <input type="text" value="0"/> Disable replication? <input type="checkbox"/> Yes		Choose “Classic RAST” for the current production RAST, or “RASTik” for the new modular RAST pipeline currently in testing. Please select which type of gene calling you would like RAST to perform. Note that using GLIMMER-3 will disable automatic error fixing, frameshift correction and the backfilling of gaps. Choose the version of FIGfams to be used to process this genome. The automatic annotation process may run into problems, such as gene candidates overlapping RIIAs, or genes embedded inside other genes. To automatically resolve these problems (even if that requires deleting some gene candidates), please check this box. If you wish for the pipeline to fix frameshifts, check this option. Otherwise frameshifts will not be corrected. If you wish RAST to build a metabolic model for this genome, check this option. If you wish for the pipeline to blast large gaps for missing genes, check this option. If you wish debug statements to be printed for this job, check this box. Set this to the verbosity level of choice for error messages. Even if this job is identical to a previous job, run it from scratch.
--	--	---

[Finish the upload](#)

Genome annotation using RAST may take half a day, an entire day or longer depending on server availability. Upon completion you can go to the Jobs Overview page where you can find your submitted jobs:

Progress bar color key:

- not started
- queued for computation
- in progress
- requires user input
- failed with an error
- successfully completed

Jobs you have access to :

Job ▲▼	Owner ▲▼	ID ▲▼	Name ▲▼	Num contigs ▲▼	Size (bp) ▲▼	Creation Date	Annotation Progress	Status
500222	Perdigao, Joao	1773.8519	Mycobacterium tuberculosis	385	4341263	2017-09-08 11:31:39	<div style="width: 100%; height: 10px; background-color: #00ff00;"></div>	not started

If your job is already completed you can click on *view details*, which will give you access to RAST output files. These files include, GenBank, GFF, FASTA, EMBL and even spreadsheets (Excel and tsv formats) containing the list of features, along with genome coordinates, orientation, gene ID, gene product, nucleotide and protein sequences, etc.

Job Details #473569

» [Browse annotated genome in SEED Viewer](#)

» Available downloads for this job: Genbank ▼ Download Update download files

» [Share this genome with selected users](#)

» View [Close Strains for this job](#)

» [Back to the Jobs Overview](#)

A RAST annotated PT000033 genome is available for you under the RAST_output sub-directory in Module2 directory. Open a terminal and type:

```
$ cd ./Module2/RAST_output
# Then start artemis:
$ art
```

Then open the GenBank file as you did with the Prokka annotation file. Is this annotation better?

Also, try to open the Excel spreadsheet with LibreOffice Calc in the Virtual Machine. LibreOffice Calc is an Open-Source alternative to Microsoft Excel. You can easily extract gene information from your strain using this file as well (and maybe even reading it into R!).

